

Machine Learning and Data Analytics

ME 5013- Fall 2019

Lectures 11

- Support Vector Machine



The University of Texas at San Antonio™


Adel Alaeddini, PhD

Associate Professor of Mechanical Engineering

Advanced Data Engineering Lab

adel.alaeddini@utsa.edu

$$h_{\beta}(x) = \frac{1}{1 + e^{-x\beta}}$$




If $y = 1$, we want $h_{\beta}(x) \approx 1$, $x\beta \gg 0$

If $y = 0$, we want $h_{\beta}(x) \approx 0$, $x\beta \ll 0$


Cost of example: $-(y \log h_{\hat{\beta}}(x) + (1 - y) \log(1 - h_{\hat{\beta}}(x)))$

$$= -y \log \frac{1}{1 + e^{-x\beta}} - (1 - y) \log\left(1 - \frac{1}{1 + e^{-x\beta}}\right)$$

If $y = 1$ (want $x\beta \gg 0$):



If $y = 0$ (want $x\beta \ll 0$):



Support vector machine

Logistic regression:

$$\min_{\beta} \frac{1}{n} \left[\sum_{i=1}^n y^{(i)} \left(-\log h_{\beta}(x^{(i)}) \right) + (1 - y^{(i)}) \left(-\log(1 - h_{\beta}(x^{(i)})) \right) \right] + \frac{\lambda}{2n} \sum_{j=1}^p \beta_j^2$$

Regularization term

Support vector machine:

$$\min_{\beta} C \sum_{i=1}^n \left[y^{(i)} \text{cost}_1(x^{(i)}\beta) + (1 - y^{(i)}) \text{cost}_0(x^{(i)}\beta) \right] + \frac{1}{2} \sum_{j=1}^p \beta_j^2$$

Regularization term



Hypothesis:

$$h_{\beta}(x) = \begin{cases} 1 & \text{if } x^{(i)}\beta \geq 0 \\ 0 & \text{Otherwise} \end{cases}$$

- No division by n
- C vs λ

Support Vector Machine

$$\min_{\beta} C \sum_{i=1}^n \left[y^{(i)} \text{cost}_1(x^{(i)}\beta) + (1 - y^{(i)}) \text{cost}_0(x^{(i)}\beta) \right] + \frac{1}{2} \sum_{j=1}^p \beta_j^2$$



If $y = 1$, we want $x\beta \geq 1$ (not just ≥ 0)


If $y = 0$, we want $x\beta \leq -1$ (not just < 0)

(extra) safety margin factor

$$\min_{\beta} C \sum_{i=1}^n \left[y^{(i)} \text{cost}_1(x^{(i)} \beta) + (1 - y^{(i)}) \text{cost}_0(x^{(i)} \beta) \right] + \frac{1}{2} \sum_{j=1}^p \beta_j^2$$


Whenever $y^{(i)} = 1$:

$$x^{(i)} \beta \geq 1$$



Whenever $y^{(i)} = 0$:

$$x^{(i)} \beta \leq -1$$




For very large C, i.e. $C=10000$

$$\min \frac{1}{2} \sum_{j=1}^p \beta_j^2$$

s. t.


$$\begin{aligned} x^{(i)} \beta &\geq 1 && \text{if } y^{(i)} = 1 \\ x^{(i)} \beta &\leq -1 && \text{if } y^{(i)} = 0 \end{aligned}$$

Linearly separable case



Large margin classifier

Large margin classifier in presence of outliers



Given $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$,
choose $l^{(1)} = x^{(1)}, l^{(2)} = x^{(2)}, \dots, l^{(m)} = x^{(m)}$.

Given example x :

$$f_1 = \text{similarity}(x, l^{(1)})$$

$$f_2 = \text{similarity}(x, l^{(2)})$$

Hypothesis: Given x , compute features $f \in \mathbb{R}^{m+1}$

Predict “y=1” if $f^{(i)}\beta \geq 0$

Training:

$$\min_{\beta} C \sum_{i=1}^n \left[y^{(i)} \text{cost}_1(f^{(i)}\beta) + (1 - y^{(i)}) \text{cost}_0(f^{(i)}\beta) \right] + \frac{1}{2} \sum_{j=1}^n \beta_j^2$$

SVM parameters:

Large C : Lower bias, high variance


Small C : Higher bias, low variance.

Large σ^2 : Features f_i vary more smoothly.

Higher bias, lower variance.

Small σ^2 : Features f_i vary less smoothly.

Lower bias, higher variance.



Use SVM software package to solve for parameters β .

Need to specify:

Choice of parameter C.

Choice of kernel (similarity function):

Linear kernel (no kernel)

Predict “y = 1” if $x\beta \geq 0$

- Large number of features
- small training set

Gaussian kernel:

$f_i = \exp\left(-\frac{\|x - l^{(i)}\|^2}{2\sigma^2}\right)$, where $l^{(i)} = x^{(i)}$. • Small number of features
Need to choose σ^2 . • Large training set

Kernel (similarity) functions:

```
function f = kernel(x1,x2)
```

$$f = \exp \left(-\frac{\| \mathbf{x1} - \mathbf{x2} \|^2}{2\sigma^2} \right)$$

```
return
```


Note: Do perform feature scaling before using the Gaussian kernel.

Other choices of kernel

Note: Not all similarity functions $\text{similarity}(x, l)$ make valid kernels. (Need to satisfy technical condition called “Mercer’s Theorem” to make sure SVM packages’ optimizations run correctly, and do not diverge).

Many off-the-shelf kernels available:

- Polynomial kernel: $(xl + \text{constant})^2$
- More esoteric: String kernel, chi-square kernel, histogram intersection kernel, ...



$$y \in \{1, 2, 3, \dots, k\}$$

Many SVM packages already have built-in multi-class classification functionality.

Otherwise, use one-vs.-all method. (Train k SVMs, one to distinguish $y = i$ from the rest, for $i = 1, \dots, k$), get $\beta^{(1)}, \dots, \beta^{(k)}$
Pick class l with largest $x\beta^{(l)}$

p = number of features ($x \in \mathbb{R}^{n+1}$), n = number of training examples

If p is large (relative to n) (i.e. $p=10,000$, $n<1000$) :

Use logistic regression, or SVM without a kernel (“linear kernel”)

If p is small (1-1000), n is intermediate (10-10,000):

Use SVM with Gaussian kernel

If p is small (1-1000), n is large ($>100,000$):

(manually) create/add more features, then use logistic regression or SVM without a kernel

Neural network likely to work well for most of these settings, but may be slower to train.