



Lecture 19

Convolutional Neural Networks

Research Questions: Visual Recognition

Research Questions: Visual Recognition

Emotion Detection



Image Source: <https://www.cbronline.com/news/automatic-facial-recognition>

Research Questions: Visual Recognition

Emotion Detection



Image Source: <https://www.cbronline.com/news/automatic-facial-recognition>

Object Detection




Image Source: <https://martechtoday.com/how-visual-recognition-is-set-to-change-advertising-223719>

Research Questions: Visual Recognition

Emotion Detection



Image Source: <https://www.cbronline.com/news/automatic-facial-recognition>

Object Detection




Image Source: <https://martechtoday.com/how-visual-recognition-is-set-to-change-advertising-223719>

Image Segmentation



Image Source: <https://www.vox.com/future-perfect/2019/4/19/18412674/ai-bias-facial-recognition-black-gay-transgender>

Research Questions: Visual Recognition

Emotion Detection



Image Source: <https://www.cbronline.com/news/automatic-facial-recognition>

Object Detection




Image Source: <https://martechtoday.com/how-visual-recognition-is-set-to-change-advertising-223719>

Image Segmentation




Image Source: <https://www.vox.com/future-perfect/2019/4/19/18412674/ai-bias-facial-recognition-black-gay-transgender>

Image Captioning



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."

Image Source: <https://towardsdatascience.com/image-captioning-in-deep-learning-9cd23fb4d8d2>

Research Questions: Visual Recognition

Emotion Detection




Image Source: <https://www.cbronline.com/news/automatic-facial-recognition>

Object Detection




Image Source: <https://martechtoday.com/how-visual-recognition-is-set-to-change-advertising-223719>

Image Segmentation



Image Source: <https://www.vox.com/future-perfect/2019/4/19/18412674/ai-bias-facial-recognition-black-gay-transgender>

Image Captioning



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."

Image Source: <https://towardsdatascience.com/image-captioning-in-deep-learning-9cd23fb4d8d2>

Pose Detection





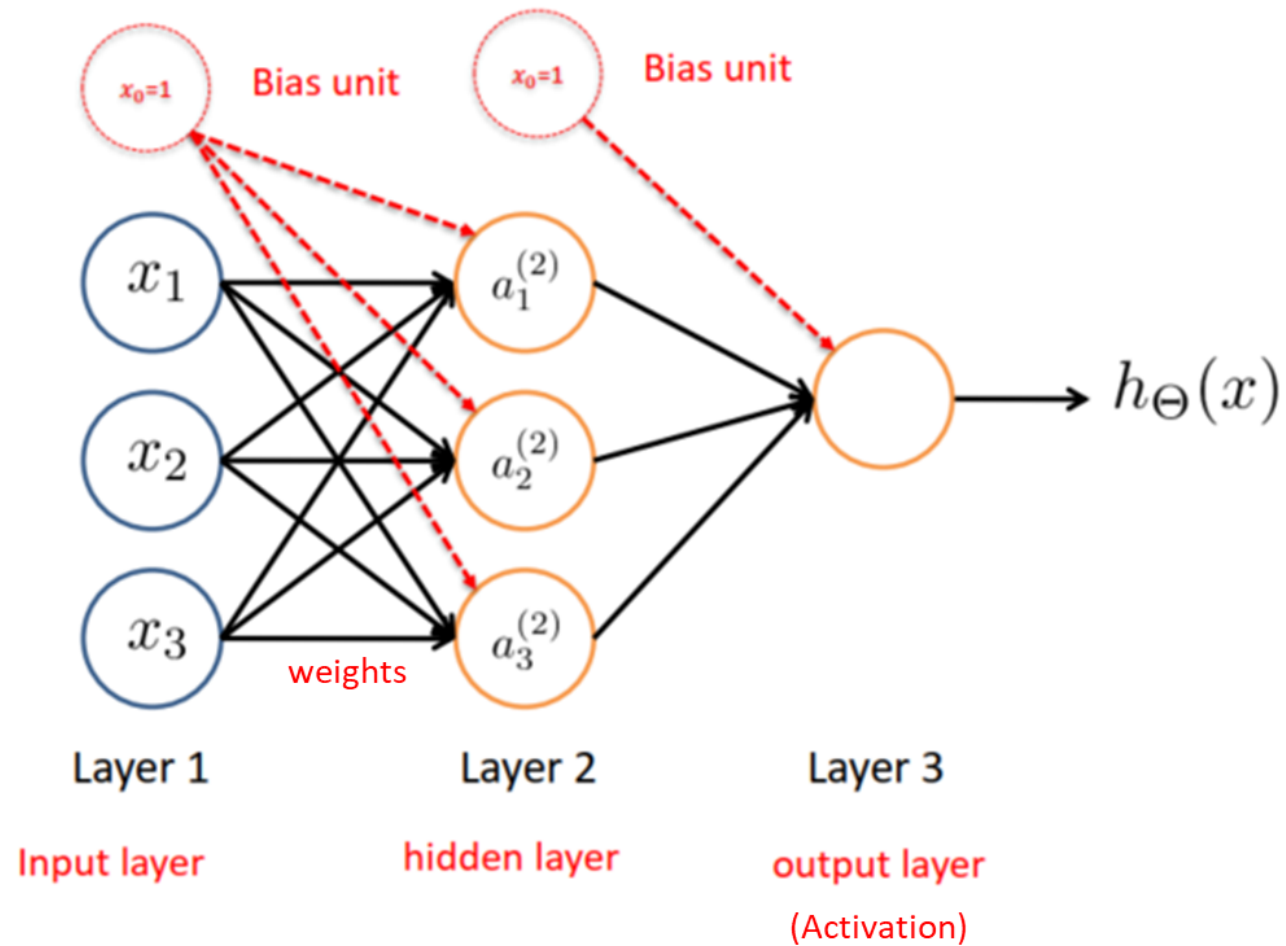
Image Source: <https://nanonets.com/blog/human-pose-estimation-2d-guide/>

Research Questions: Image Generation




Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019.

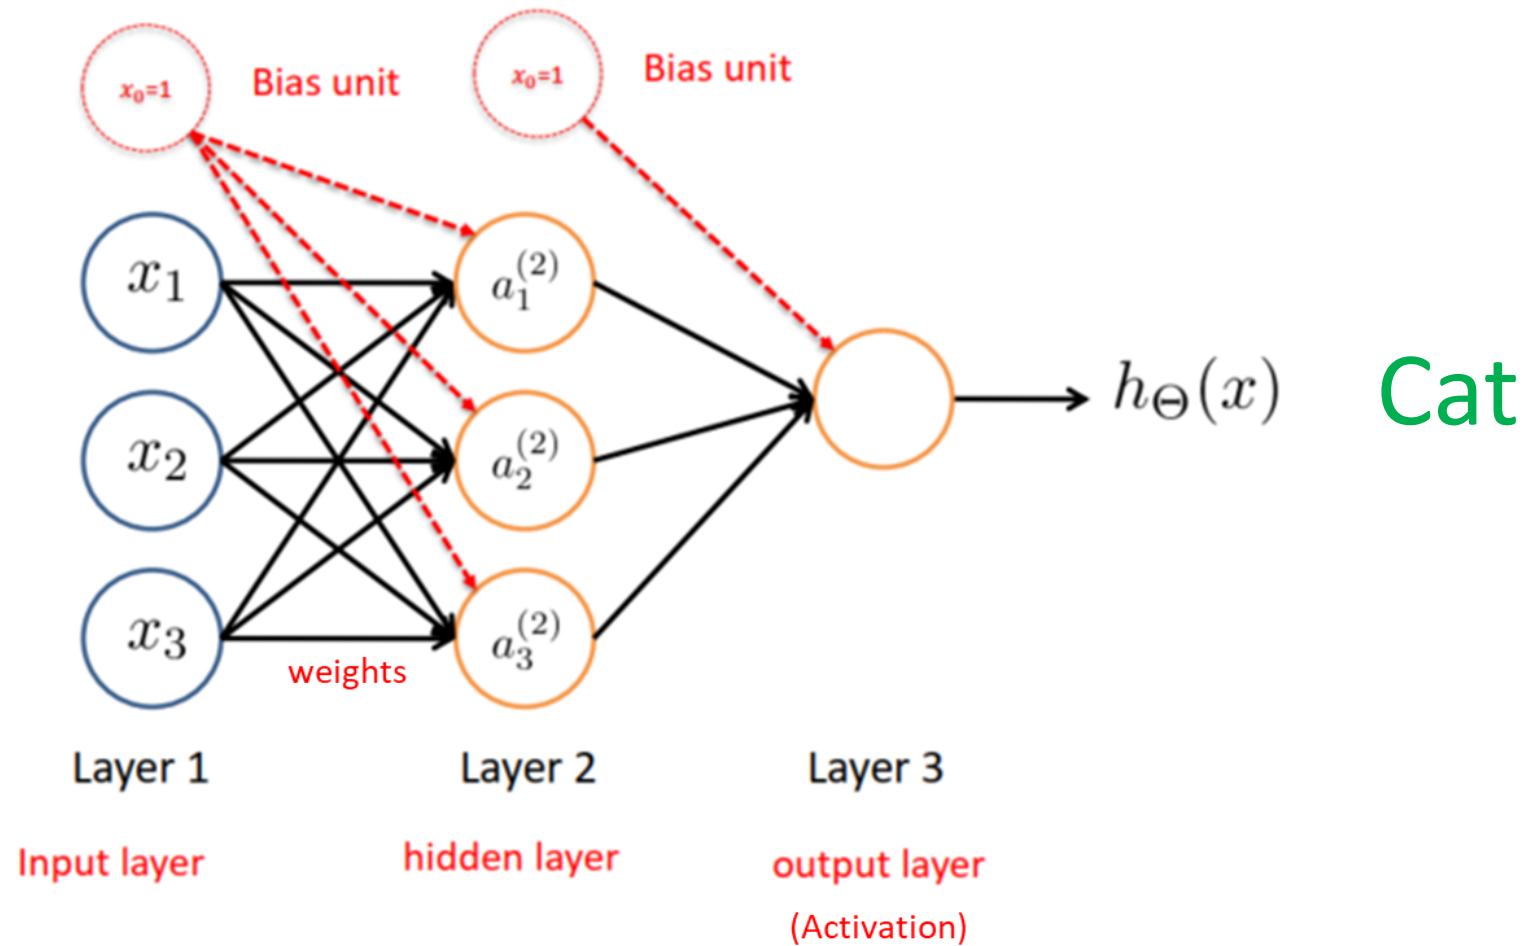
Recall from last time



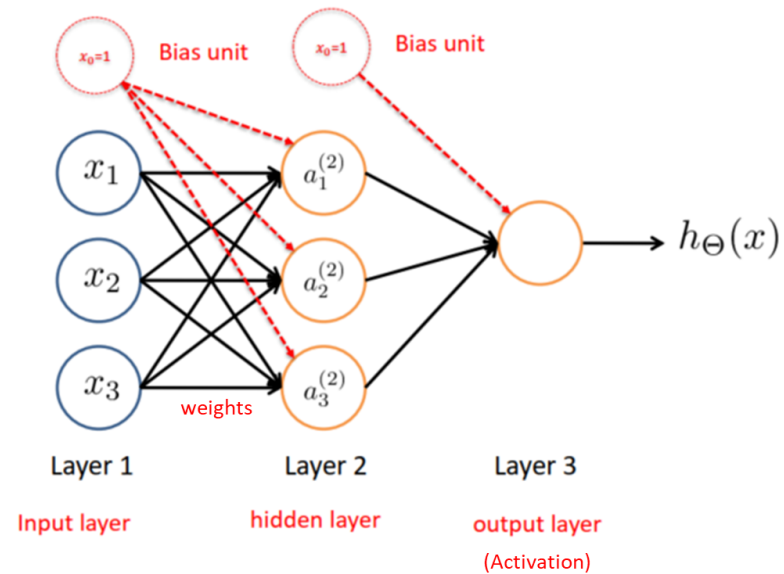
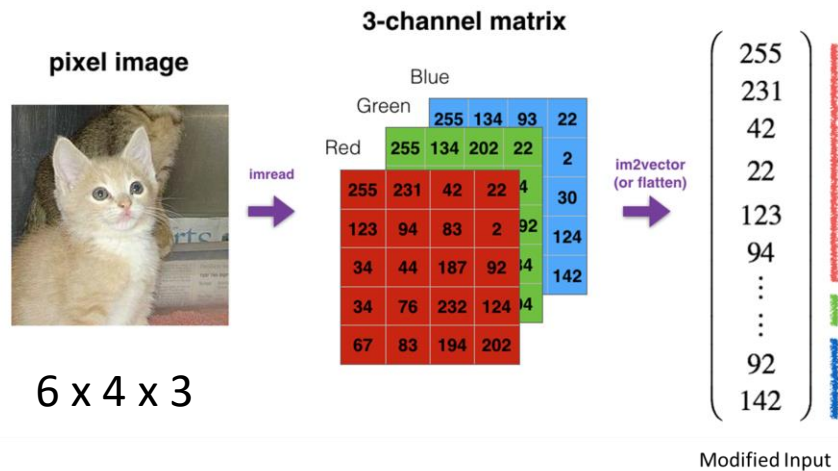
Recall from last time




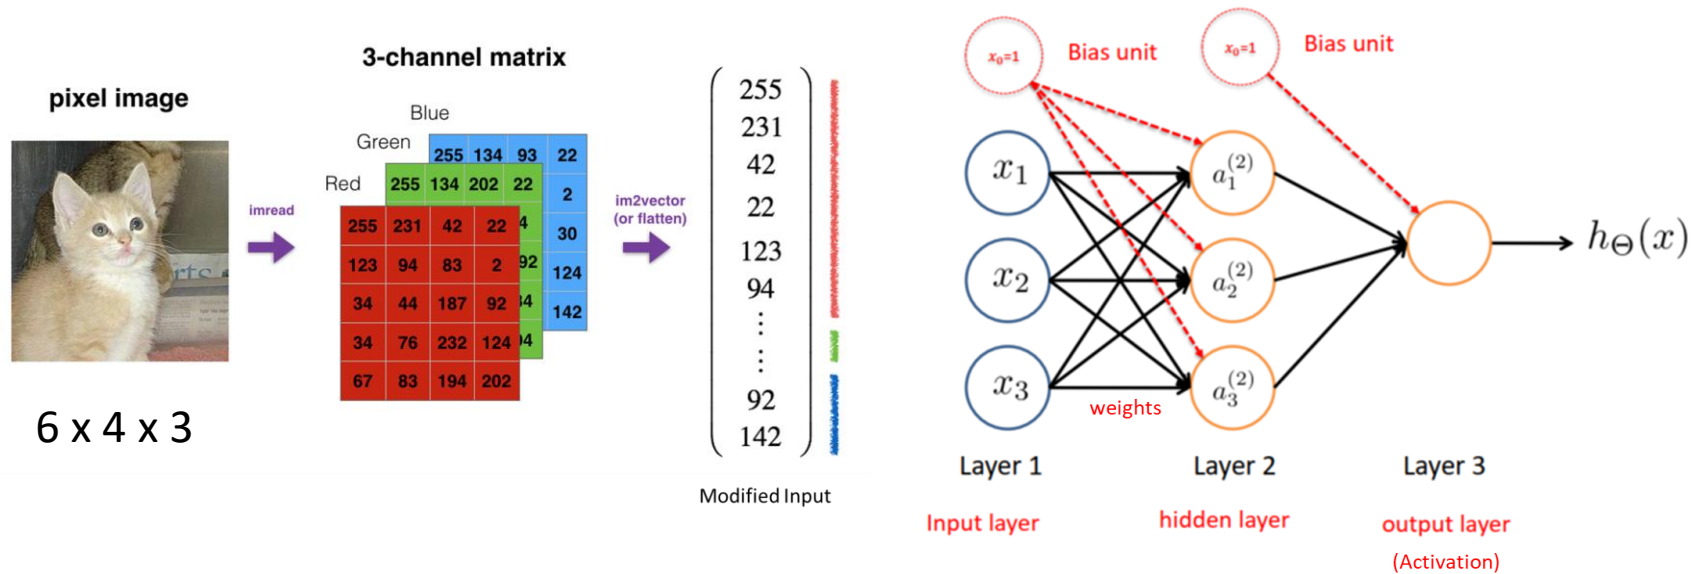
Recall from last time



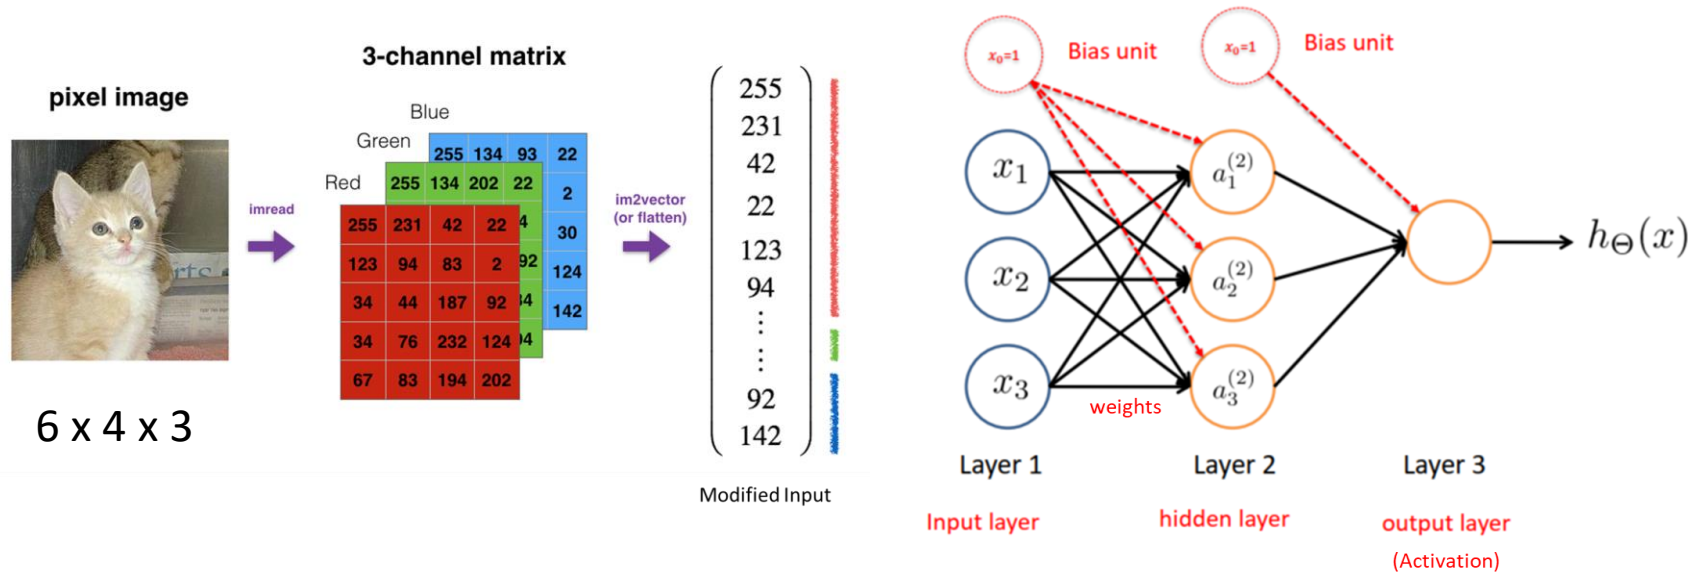
Fully Connected Layers



Fully Connected Layers





Fully Connected Layers


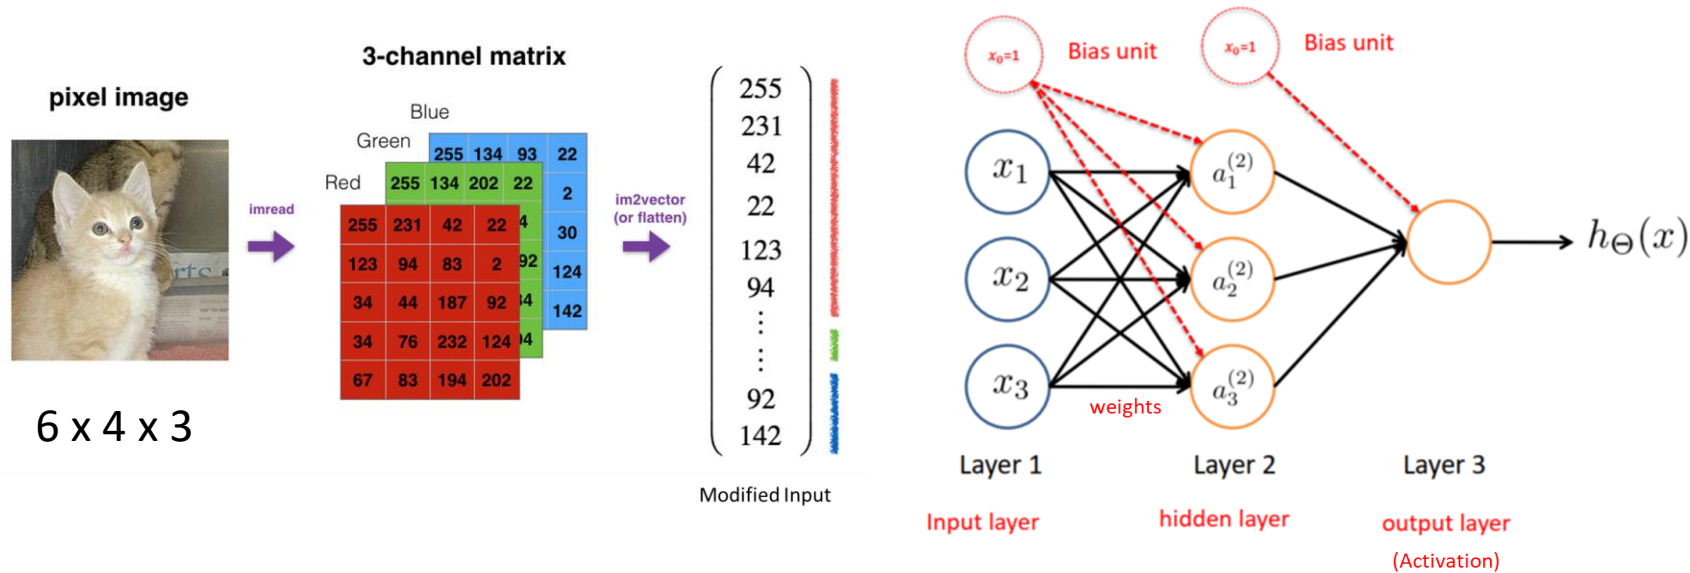


$$\begin{array}{c} \text{Input} \\ 72 \end{array} \longrightarrow \begin{array}{c} Wx \\ 2 \times 72 \end{array}$$

Fully Connected Layers



Fully Connected Layers



Output: 1 Number

The result of taking a dot product between a row of W and the input (a 72-dimensional dot product)

Issues in Image Recognition using Fully Connected Networks

Issues in Image Recognition using Fully Connected Networks

1. They don't scale well to full images



Issues in Image Recognition using Fully Connected Networks

1. *They don't scale well to full images*



$32 \times 32 \times 3 = 3072$ weights

Width x Height x Color Channel

Issues in Image Recognition using Fully Connected Networks

1. *They don't scale well to full images*



$32 \times 32 \times 3 = 3072$ weights

Width x Height x Color Channel

What if we have a color image of $500 \times 500 \times 3$
(750000 Weights)

For multiple layers, number of weights to calculate will Increase!

Issues in Image Recognition using Fully Connected Networks

1. *They don't scale well to full images*



$32 \times 32 \times 3 = 3072$ weights

Width x Height x Color Channel

What if we have a color image of $500 \times 500 \times 3$
(750000 Weights)

For multiple layers, number of weights to calculate will Increase!

2. *Doesn't preserve the spatial structure*




Convolution Neural Networks

- We use three main types of layers to build Convolution Neural Network architectures
 1. **Convolutional Layer** (combined with activation function),
 2. **Pooling Layer**, and
 3. **Fully-Connected Layer (Exactly NN)**


Convolution Layer (Higher Level Overview)

Convolution Layer (Higher Level Overview)




Input Image
32 x 32 x 3


Convolution Layer (Higher Level Overview)




Convolution Layer (Higher Level Overview)



Convolution Layer (Higher Level Overview)



Convolution Layer (Higher Level Overview)



Input Image
32 x 32 x 3



Filter
5 x 5 x 3


1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved
Feature


Convolution Layer (Higher Level Overview)



Convolution Layer (Higher Level Overview)



Convolution Layer (Higher Level Overview)

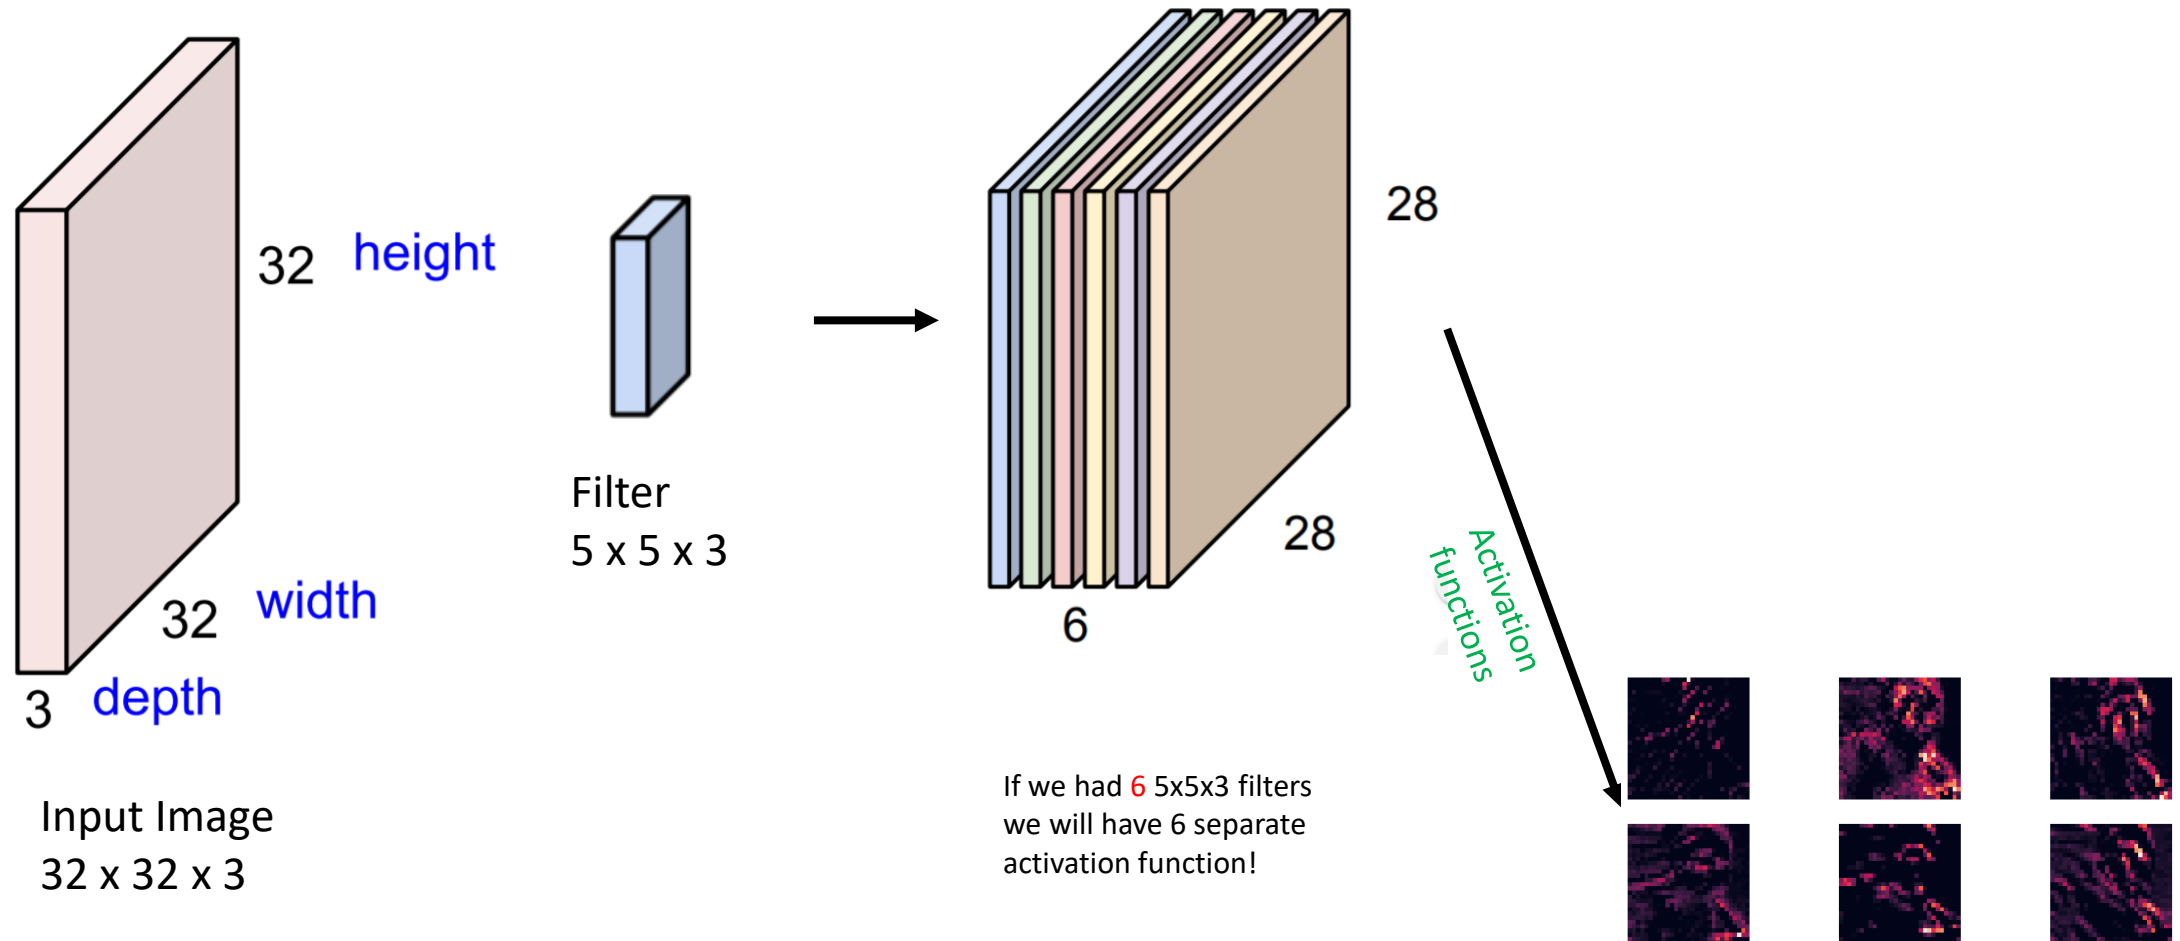


Input Image
32 x 32 x 3

Filter
5 x 5 x 3

If we had 6 5x5x3 filters
we will have 6 separate
activation function!

Convolution Layer (Higher Level Overview)



Convolution Layer (Filter)

Convolution Layer (Filter)

How is it helping?



*

1	0	-1
2	0	-2
1	0	-1



Convolution Layer (Filter)

How is it helping?




*

1	0	-1
2	0	-2
1	0	-1



Let's do an Example



Multiply the window element by element with the flip(kernel), then sum the results



Convolution Layer (Filter)

How is it helping?




*

1	0	-1
2	0	-2
1	0	-1



Let's do an Example



Multiply the window element by element with the flip(kernel), then sum the results

1	3	1
0	-1	1
2	2	-1

$$\Rightarrow 1(-1) + 3(0) + 0(2) - 1(1) = -2$$

1	3	1
0	-1	1
2	2	-1

$$\Rightarrow 3(-1) + 1(0) - 1(2) + 1(1) = -4$$

1	3	1
0	-1	1
2	2	-1

$$\Rightarrow 0(-1) - 1(0) + 2(2) + 2(1) = 6$$

1	3	1
0	-1	1
2	2	-1

$$\Rightarrow -1(-1) + 1(0) + 2(2) - 1(1) = 4$$

result(valid)


-2	-4
6	4

Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.


Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.
- Do you think we can change the strides?



Convolution Layer



- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.
- Do you think we can change the strides? **Yes!**



Stride : 1


Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.
- Do you think we can change the strides? **Yes!**




Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.
- Do you think we can change the strides? **Yes!**




Stride : 1




Stride : 2

Convolution Layer

- While convoluting so far, we have shown a sliding window of 1. This is called a ***STRIDE***.
- Do you think we can change the strides? **Yes!**




Stride : 1



Stride : 2

Convolution Layer

- What is the dimension of the transformed image?




Input Dimension (N): 7x7 (spatially)

Filter Dimension (F) : 3x3 (assume)

Stride : 1

Convolution Layer

- What is the dimension of the transformed image?




Input Dimension (N): 7x7 (spatially)
Filter Dimension (F): 3x3 (assume)
Stride: 1

Output Dimension: $\frac{N - F}{Stride} + 1$

Convolution Layer

- What is the dimension of the transformed image?



Input Dimension (N): 7x7 (spatially)

Filter Dimension (F) : 3x3 (assume)


Stride : 1

Output Dimension: 5x5 (output)

$$\text{Output Dimension: } \frac{N - F}{\text{Stride}} + 1$$

Convolution Layer

- What is the dimension of the transformed image?




Input Dimension (N): 7x7 (spatially)
Filter Dimension (F) : 3x3 (assume)
Stride : 2

Output Dimension: $\frac{N - F}{Stride} + 1$

Convolution Layer

- What is the dimension of the transformed image?



Input Dimension (N): 7x7 (spatially)

Filter Dimension (F) : 3x3 (assume)


Stride : 2

Output Dimension: 3x3 (output)

$$\text{Output Dimension: } \frac{N - F}{\text{Stride}} + 1$$

Convolution Layer

- What is the dimension of the transformed image?



Input Dimension (N): 7x7 (spatially)

Filter Dimension (F): 3x3 (assume)


Stride : 3

Output Dimension: ?

$$\text{Output Dimension: } \frac{N - F}{\text{Stride}} + 1$$

Convolution Layer

- What is the dimension of the transformed image?



Input Dimension (N): 7x7 (spatially)

Filter Dimension (F): 3x3 (assume)

Stride : 3


Output Dimension: ?

$$\text{Output Dimension: } \frac{N - F}{\text{Stride}} + 1$$

Doesn't fit!

Convolution Layer

- What is the dimension of the transformed image?



Convolution Layer

- What is the dimension of the transformed image?

In practice: Common to zero pad (Z) the border!

0	0	0	0	0	0			
0								
0								
0								
0								

Pad with 1 pixel


Convolution Layer

- What is the dimension of the transformed image?

In practice: Common to zero pad (Z) the border!

0	0	0	0	0	0			
0								
0								
0								
0								

Pad with 1 pixel



Padding (P): 1


Convolution Layer

- What is the dimension of the transformed image?

In practice: Common to zero pad (Z) the border!

0	0	0	0	0	0			
0								
0								
0								
0								

Pad with 1 pixel




Padding (P): 1


Convolution Layer

- What is the dimension of the transformed image?

In practice: Common to zero pad (Z) the border!



Pad with 1 pixel




$$\text{Output Dimension: } \frac{N - F + 2P}{\text{Stride}} + 1$$

Padding (P): 1


Convolution Layer

- What is the dimension of the transformed image?

In practice: Common to zero pad (Z) the border!



Pad with 1 pixel



$$\text{Output Dimension: } \frac{N - F + 2P}{\text{Stride}} + 1$$

Padding (P): 1

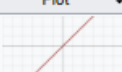
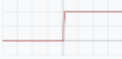



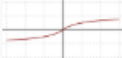
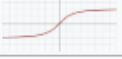







Output Dimension: 3 x 3

Convolution Layer

- In summary, we have **four** important hyper-parameters:
 1. Number of Filters, K
 2. Size of the Filter, F
 3. Stride, S
 4. Zero Padding, Z

Activation Function

- A value passed through the function will be transformed within a range.


Name	Plot	Equation
Identity		
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Sigmoid or Soft step)		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}^{[1]}$
TanH		$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$
SQLN ^[9]		$f(x) = \begin{cases} 1 & : x > 2.0 \\ x - \frac{x^2}{4} & : 0 \leq x \leq 2.0 \\ x + \frac{x^2}{4} & : -2.0 \leq x < 0 \\ -1 & : x < -2.0 \end{cases}$
ArcTan		$f(x) = \tan^{-1}(x)$
ArSinH		$f(x) = \sinh^{-1}(x) = \ln(x + \sqrt{x^2 + 1})$
ElliotSig ^{[10][11]} Softsign ^{[12][13]}		$f(x) = \frac{x}{1 + x }$
Inverse square root unit (ISRU) ^[14]		$f(x) = \frac{x}{\sqrt{1 + \alpha x^2}}$
Inverse square root linear unit (ISRLU) ^[14]		$f(x) = \begin{cases} \frac{x}{\sqrt{1 + \alpha x^2}} & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Rectified linear unit (ReLU) ^[15]		$f(x) = \begin{cases} 0 & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases}$
Bipolar rectified linear unit (BReLU) ^[16]		$f(x_i) = \begin{cases} ReLU(x_i) & \text{if } i \bmod 2 = 0 \\ -ReLU(-x_i) & \text{if } i \bmod 2 \neq 0 \end{cases}$
Leaky rectified linear unit (Leaky ReLU) ^[17]		$f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Parameteric rectified linear unit (PReLU) ^[18]		$f(\alpha, x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Randomized leaky rectified linear unit (RReLU) ^[19]		$f(\alpha, x) = \begin{cases} \alpha x & \text{for } x < 0^{[3]} \\ x & \text{for } x \geq 0 \end{cases}$
Exponential linear unit (ELU) ^[20]		$f(\alpha, x) = \begin{cases} \alpha(e^x - 1) & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases}$
Scaled exponential linear unit (SELU) ^[21]		$f(\alpha, x) = \lambda \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$ with $\lambda = 1.0507$ and $\alpha = 1.67326$
S-shaped rectified linear activation unit (SReLU) ^[22]		$f_{t_l, t_r, a_l, a_r}(x) = \begin{cases} t_l + a_l(x - t_l) & \text{for } x \leq t_l \\ x & \text{for } t_l < x < t_r \\ t_r + a_r(x - t_r) & \text{for } x \geq t_r \end{cases}$ t_l, a_l, t_r, a_r are parameters.
Adaptive piecewise linear (APL) ^[23]		$f(x) = \max(0, x) + \sum_{i=1}^S \alpha_i^+ \max(0, -x + b_i^+)$

Activation Function

- A value passed through the function will be transformed within a range.


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Activation Function

- A value passed through the function will be transformed within a range.
- The most used activation function in CNNs is the ReLU (Rectified Linear Unit).


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Activation Function

- A value passed through the function will be transformed within a range.
- The most used activation function in CNNs is the ReLU (Rectified Linear Unit).


1	14	-9	4
-2	-20	10	6
-3	3	11	1
2	54	-2	80

ReLU

1	14	0	4
0	0	10	6
0	3	11	1
2	54	0	80


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Activation Function

- A value passed through the function will be transformed within a range.
- The most used activation function in CNNs is the ReLU (Rectified Linear Unit).


1	14	-9	4
-2	-20	10	6
-3	3	11	1
2	54	-2	80

ReLU


1	14	0	4
0	0	10	6
0	3	11	1
2	54	0	80

- While using **ReLU**, be careful with your selection of learning rate!


Sigmoid
 $\sigma(x) = \frac{1}{1+e^{-x}}$




tanh
 $\tanh(x)$



ReLU
 $\max(0, x)$




Leaky ReLU
 $\max(0.1x, x)$




Maxout
 $\max(w_1^T x + b_1, w_2^T x + b_2)$


ELU
 $\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$



Convolution + Activation Layer




Convolution + Activation Layer



Pooling Layers

- Pooling in a CNN is a subsampling step
 - Replaces the output to summary statistics of nearby outputs!
- It is common to periodically insert a Pooling layer in-between successive Conv layers
- Operates over each activation map independently
- Pooling is performed for two reasons
 1. Dimensionality reduction.
 2. Invariance to transformations of rotation and translation (*If we translate the input by a small amount values of most of the outputs does not change*). This is important to identify if a feature is present rather than where exactly it is.



Pooling Layers

Similar to Convolution layer the strides can be changed!

0	1	0	4	5
2	3	2	1	3
4	4	0	4	3
2	5	2	6	4
1	0	0	5	7

4	4	5
5	6	6
5	6	7

Max Pooling
3x3, Stride 1

1. norm of a rectangular neighborhood
2. Weighted average based on the distance from the central pixel

Pooling Layers

Similar to Convolution layer the strides can be changed!

0	1	0	4	5
2	3	2	1	3
4	4	0	4	3
2	5	2	6	4
1	0	0	5	7

4	4	5
5	6	6
5	6	7

Max Pooling
3x3, Stride 1

4	5
5	6

Max Pooling
3x3, Stride 2

1. norm of a rectangular neighborhood
2. Weighted average based on the distance from the central pixel

Pooling Layers

L_2 norm of a rectangular neighborhood

Similar to Convolution layer the strides can be changed!

- Popular pooling functions:
 1. Max pooling operation reports the maximum output within a rectangular neighborhood
 2. Average of a rectangular neighborhood
 3. Sum pooling
 4. Weighted average based on the distance from the central pixel

0	1	0	4	5
2	3	2	1	3
4	4	0	4	3
2	5	2	6	4
1	0	0	5	7

4	4	5
5	6	6
5	6	7

Max Pooling
3x3, Stride 1

4	5
5	6

Max Pooling
3x3, Stride 2

Feature Map

6	6	6	6
4	5	5	4
2	4	4	2
2	4	4	2

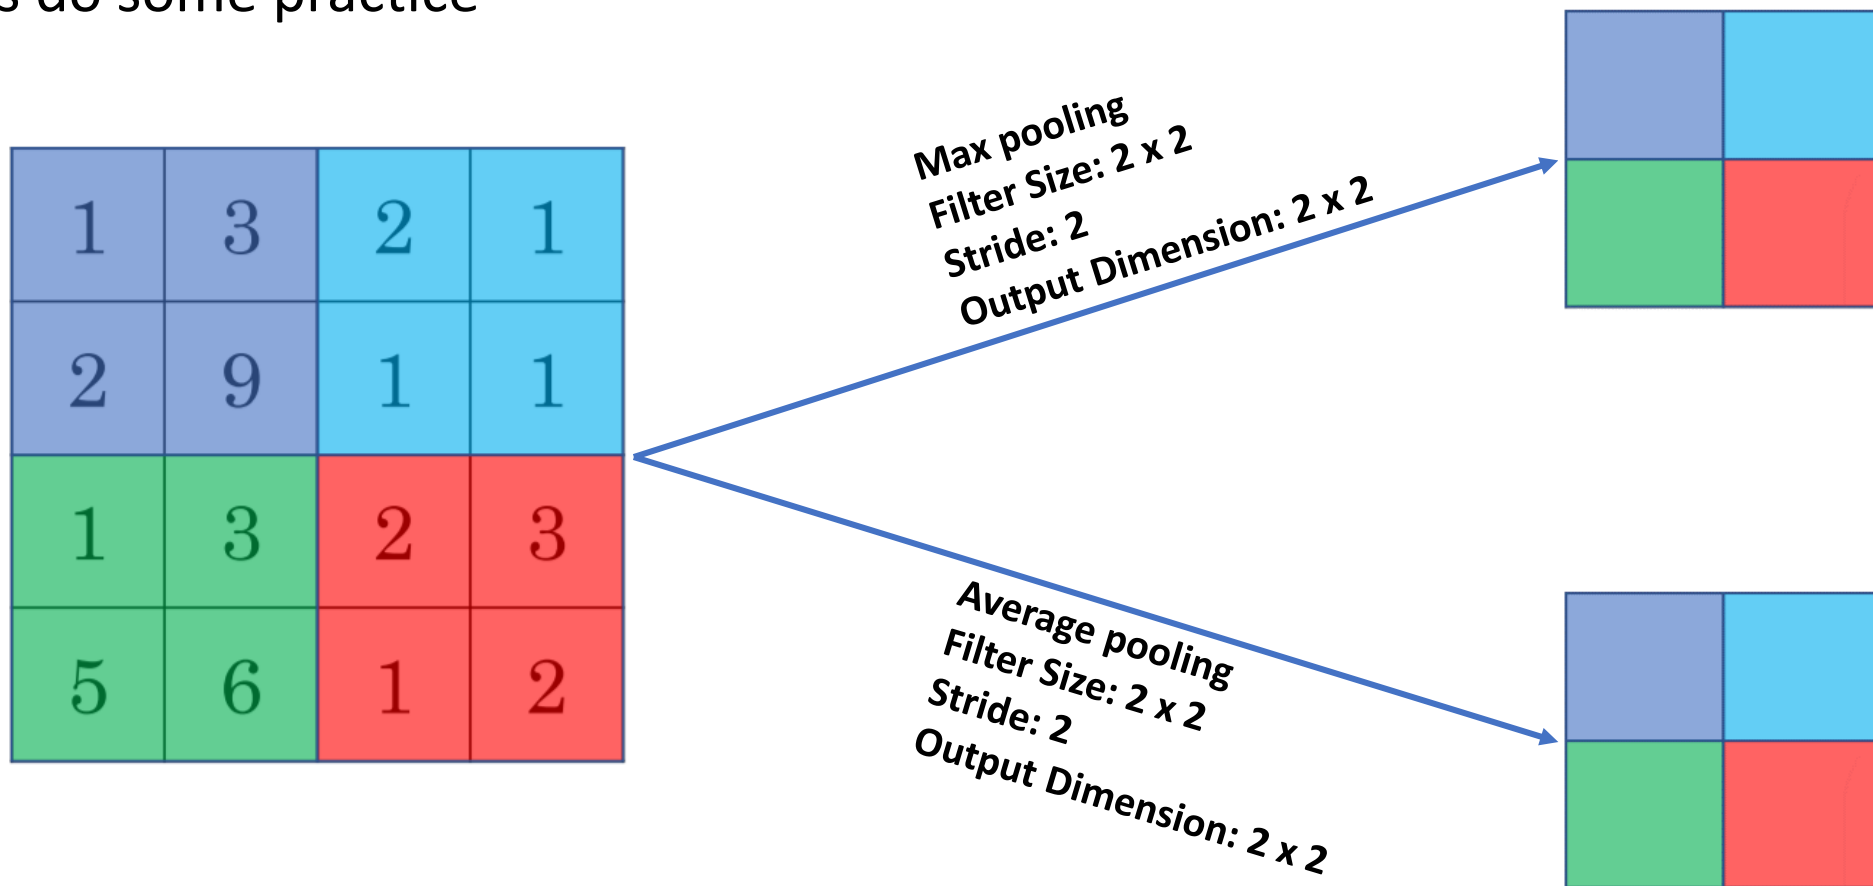
Max
Pooling

Average
Pooling

Sum
Pooling

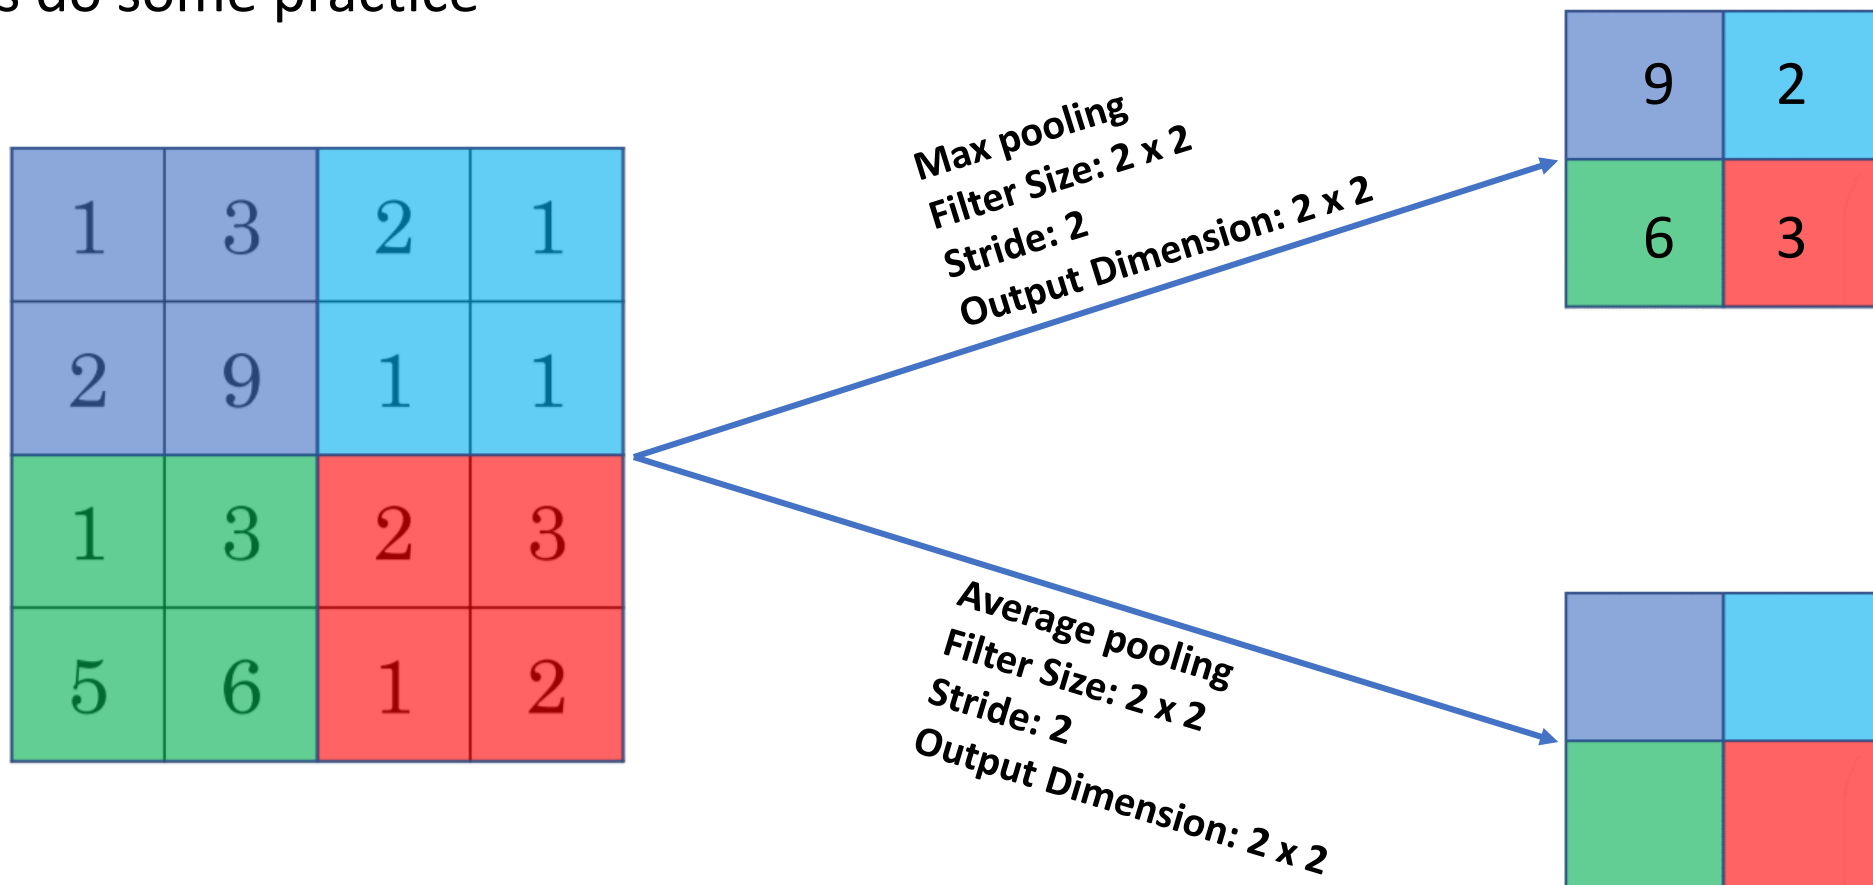
Pooling Layers

- Let's do some practice



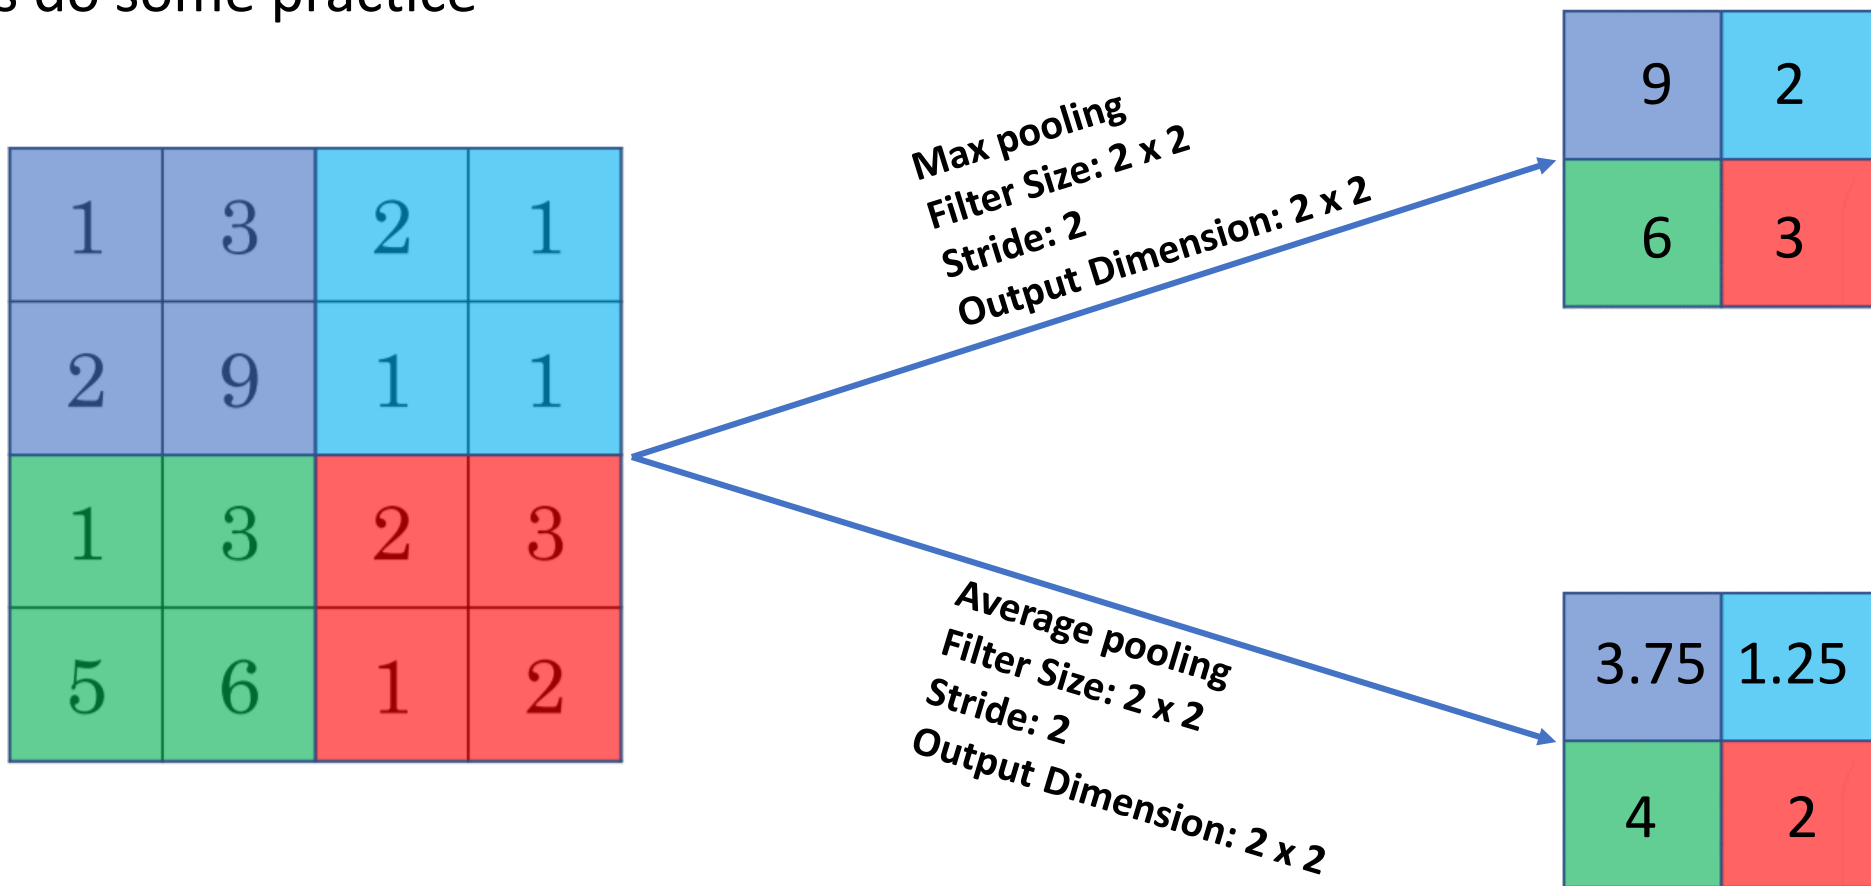
Pooling Layers

- Let's do some practice



Pooling Layers

- Let's do some practice

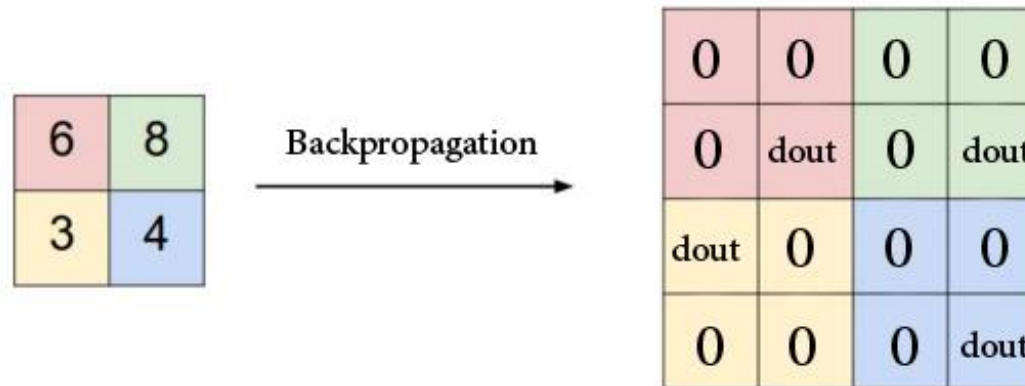


Pooling Layers

- In summary, we have **three** important hyper-parameters:
 1. Filters Size, F
 2. Stride, S
 3. Zero Padding are rarely used
- Max pooling or Average pooling popularly are used


Un-Pooling Layers

- While backpropagating we only preserve the max nodes:



Example 1 (Classic LeNet – 5)

- This model was published in 1998. The last layer wasn't using softmax back then.
- Used for digit recognition.
- The dimensions of the image decreases as the number of channels increases.
- It has *60k* parameters.



Example 2

