

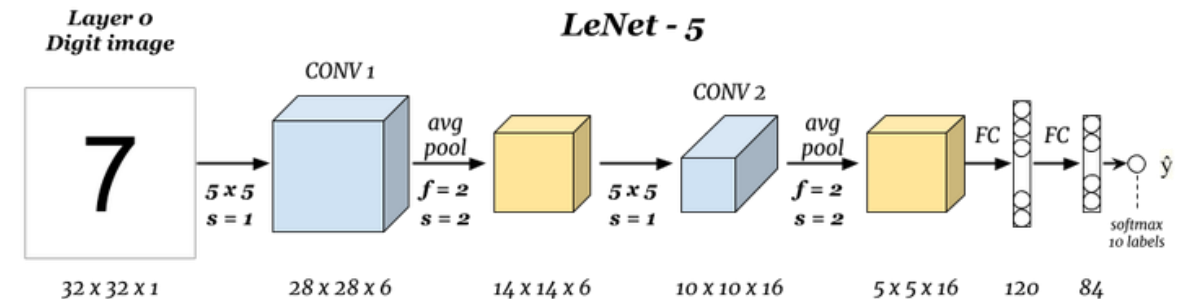
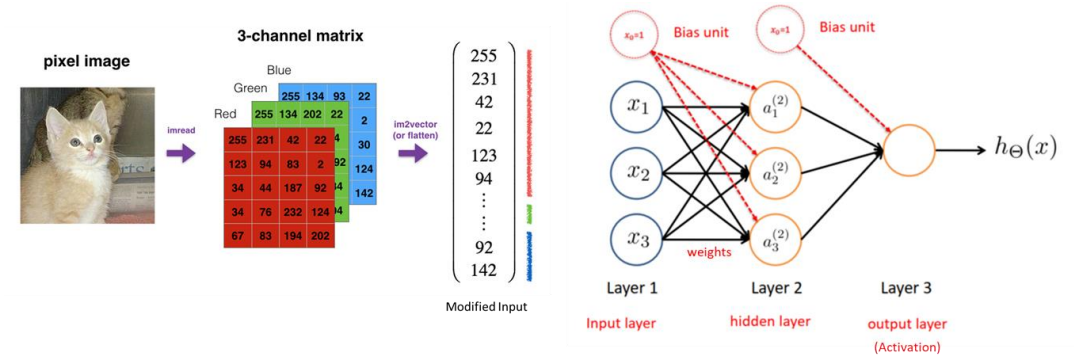


Lecture 20

Parameter Learning in NN



Quick recap on what we learned so far

- Fully Neural Network
- Convolutional Neural Network
- An NN architecture
 - How many layers?
 - How many neurons? and
 - How the neurons are connected?
- The parameters values (weights!)
- Components
 - Filters:
 - Convolutional Filter
 - Pooling Filters
 - Activation Functions
 - Sigmoid
 - ReLu etc.

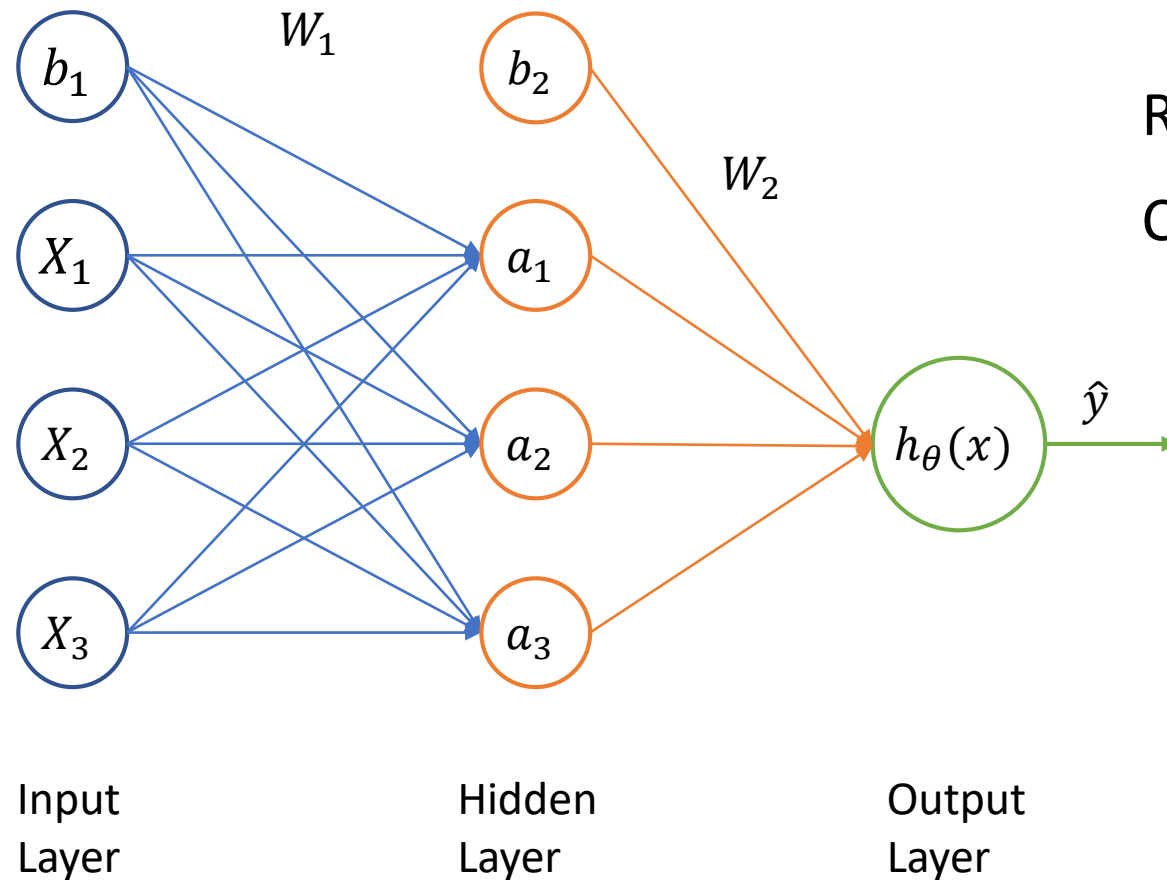


Quick recap on what we learned so far

- How many layers should we use?
 - Theoretically:
 - A NN with one Hidden layer is a universal function approximator (*Cybenko, 1989*)
 - Empirically:
 - Before year 2006 : There should be at least 2 – 6 hidden layers.
 - After year 2006 : There should be at least 5+ hidden layers.
- How the decision boundary works?
 - 0 Hidden Layer
 - 1 Hidden Layer
 - 2 Hidden Layer



Quick recap on what we learned so far



Regression Problem, $\hat{y}_r = W_2^T a_i + b_2$


Classification Problem, $\hat{y}_c = \sigma(W_2^T a_i + b_2)$

↓

$$\sigma(x) = \frac{1}{1+e^{-x}}$$


Learning from Network

- The two Components:
 - Architecture
 - Parameters (W_1, W_2)
- Learning Steps:
 1. Initialize the **weights**, W_n
 2. Calculate the **forward propagation**
 3. Calculate the loss function
 4. Perform **backpropagation**
 5. Update the parameters
 6. Repeat until convergence!



Step 1: Weight Initialization


- **Q1.** How to initialize the **weights, W_n** ?
 - Randomly initialize the parameters to small values (e.g., normally distributed round zero; $\mathcal{N}(0, 0.1)$).
- **Q2.** What will happen if we initialize the weights as **zero**?
 - The first layer will always be the same since, $W^{[1]}x^i + b^{[1]} = 0^{3 \times 1}x^i + 0^{3 \times 1}$ where, $0^{n \times m}$ denotes a matrix of size $n \times m$ filled with zeros.
 - Again, the output of the network has a sigmoid function, thus no matter what input we provide we will get an output probability of 0.5 i.e. $\sigma(0) = 0.5$.




Step 1: Weight Initialization

- **Q3.** What if we initialized all parameters to be the same non-zero value?
 - Each element of the activation vector will be the same (because $W^{[1]}$ contains all the same values). This behavior will occur at all layers of the neural network. As a result, when we compute the gradient, all neurons in a layer will be equally responsible for anything contributed to the final loss
 - **Summary:** *All neurons will learn the same thing.*
- **Q4.** Is there a better approach than randomly initializing?
 - Yes, there is one! It's called Xavier/He initialization.

$$w^{[\ell]} \sim \mathcal{N}\left(0, \sqrt{\frac{2}{n^{[\ell]} + n^{[\ell-1]}}}\right)$$



Step 2: Forward Propagation



Regression Problem, $\hat{y}_r = W_2^T a_i + b_2$

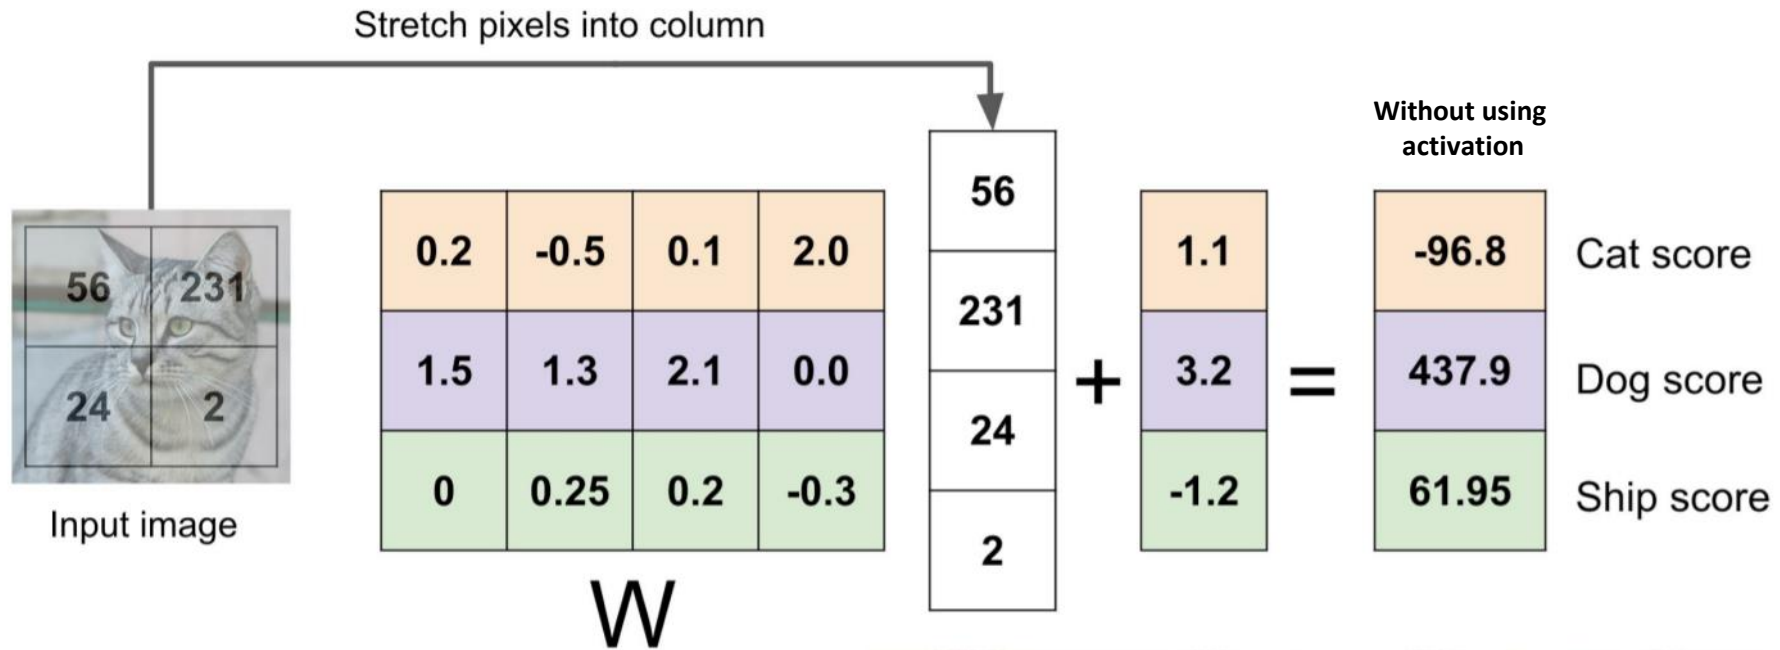
Classification Problem, $\hat{y}_c = \sigma(W_2^T a_i + b_2)$

↓

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

Hidden Layer Calculation (with activation),
 $a_i = \sigma(W_1^T X + b_1)$

Step 2: Forward Propagation



Step 3: Loss Function

- Loss function is defined as –

$$\mathcal{L}(y, \hat{y}, X) = \frac{1}{2} (y - \hat{y})^2$$

Inputs Variable

Predicted outcome

True Outcome

- In terms of log-loss, the loss function is defined as –



$$\mathcal{L}(y, \hat{y}, X) = -[(1 - y) \log(1 - \hat{y}) + y \log \hat{y}]$$

Step 4: Back Propagation


- The partial derivative of \mathcal{L} can be decomposed as a sum of partial derivatives of individual losses:

$$\frac{\partial \mathcal{L}}{\partial W^l} = \sum_{k=1}^n \frac{\partial \mathcal{L}(y, \hat{y})}{\partial W^l}$$

- Say, we have the network parameters, $W^1, W^2, W^3, b^1, b^2, b^3$.
- We will first calculate the gradient with respect to W^3 as its due to influence of W^1 on output is complex than that of W^3 .



Step 4: Updating the Parameters (Back Propagation)



- Thus we have,

$$\begin{aligned}
 \frac{\partial \mathcal{L}}{\partial W^3} &= -\frac{\partial}{\partial W^3} \left((1 - y) \log(1 - \hat{y}) + y \log \hat{y} \right) \\
 &= -(1 - y) \frac{\partial}{\partial W^3} \log(1 - \sigma(1 - W^3 a^2 + b^3)) - y \frac{\partial}{\partial W^3} \log(\sigma(W^3 a^2 + b^3)) \\
 &= (1 - y) \sigma(W^3 a^2 + b^3) a^{2T} - y (1 - \sigma(W^3 a^2 + b^3)) a^{2T} \\
 &= (1 - y) a^3 a^{2T} - y (1 - a^3) a^{2T} \\
 &= (a^3 - y) a^{2T}
 \end{aligned}$$

- To compute the gradient with respect to W^2 :

$$\frac{\partial \mathcal{L}}{\partial W^2} = \underbrace{\frac{\partial \mathcal{L}}{\partial a^3}}_{a^3 - y} \underbrace{\frac{\partial a^3}{\partial z^3}}_{W^3} \underbrace{\frac{\partial z^3}{\partial a^2}}_{\sigma'(z^2)} \underbrace{\frac{\partial a^2}{\partial z^2}}_{a^1} \frac{\partial z^2}{\partial W^2}$$

Note: $\sigma' = \sigma(1 - \sigma)$

Step 4.1: Regularization (if applicable)

- Similar to what we covered in Lecture 13. (Remember for linear regression)
- Is used to reduce the overfitting of neural networks.

Hypothesis: $h_{\beta}(x) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$

Cost Function: $J(\beta) = \frac{1}{2n} \left[\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_1^{(i)} - \dots - \beta_p x_p^{(i)})^2 + \lambda \sum_{j=1}^p \beta_j^2 \right]$

Parameters: $\beta_0, \beta_1, \dots, \beta_p$

$\min_{\beta} J(\beta)$

Gradient descent:

Repeat until convergence{

$$\beta_0 := \beta_0 + \alpha \frac{1}{n} \left[\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_1^{(i)} - \dots - \beta_p x_p^{(i)}) \right]$$

$$\beta_1 := \beta_1 + \alpha \frac{1}{n} \left[\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_1^{(i)} - \dots - \beta_p x_p^{(i)}) x_1^{(i)} - \lambda \beta_1 \right]$$

...

$$\beta_p := \beta_p + \alpha \frac{1}{n} \left[\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_1^{(i)} - \dots - \beta_p x_p^{(i)}) x_p^{(i)} - \lambda \beta_p \right]$$

}

NOTE:

$$\beta_p := \beta_p \left(1 - \alpha \frac{\lambda}{n} \right) + \alpha \frac{1}{n} \left[\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_1^{(i)} - \dots - \beta_p x_p^{(i)}) x_p^{(i)} \right]$$


Step 5: Updating the Parameters

- Optimize/update the parameters by using gradient descent optimization. For a given number of layer, l :


$$W^l = W^l - \alpha \frac{\partial \mathcal{L}}{\partial W^l}$$

$$b^l = b^l - \alpha \frac{\partial \mathcal{L}}{\partial b^l}$$

- This computation is non-trivial at hidden layers ($l < L$ (Output Layer)) but has complex chain of influence via activation values at subsequent layers.
- This problem is already addressed by **backpropagation**.




Example: Backpropagation



$$f = X_1 - X_2$$

$$\frac{\partial f}{\partial X_1} = 1 \quad \frac{\partial f}{\partial X_2} = -1$$




$$f = X_1 - X_2 = X_1 - (X_3 + X_4)$$

$$\frac{\partial f}{\partial X_1} = 1 \quad \frac{\partial f}{\partial X_2} = -1$$

$$\frac{\partial f}{\partial X_3} = ?$$

$$\frac{\partial f}{\partial X_3} = \frac{\partial f}{\partial X_2} \frac{\partial X_2}{\partial X_3} = -1 \times 1 = -1$$

Example: Backpropagation

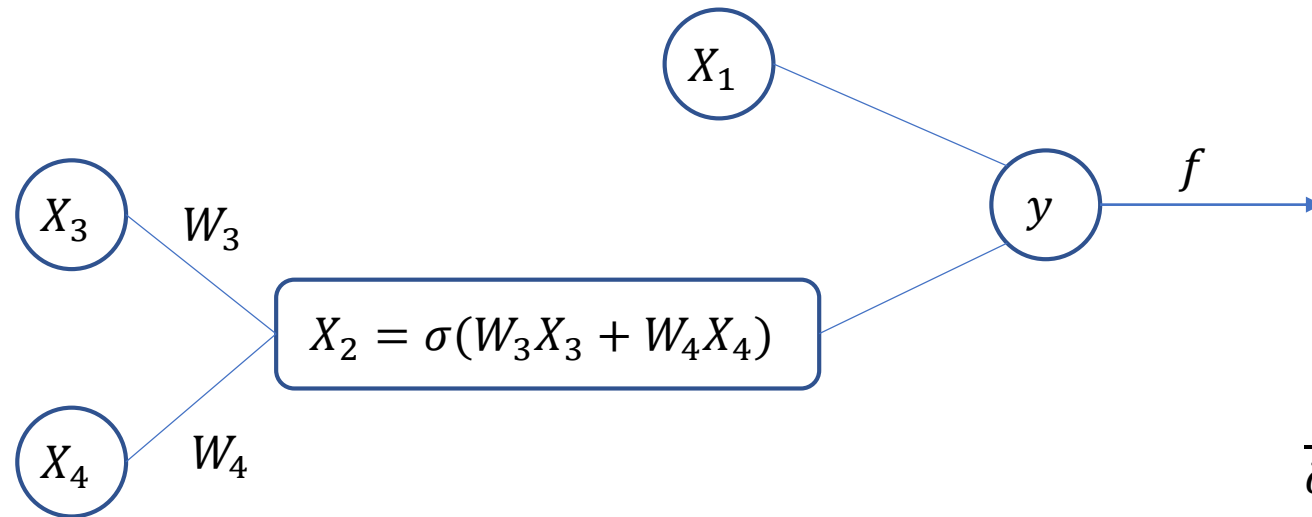


$$f = X_1 - X_2 = X_1 - (W_3X_3 + W_4X_4)$$

$$\frac{\partial X_2}{\partial X_3} = 1 \quad \frac{\partial X_2}{\partial X_4} = 1$$

$$\frac{\partial f}{\partial W_3} = ?$$

$$\frac{\partial f}{\partial W_3} = \frac{\partial f}{\partial X_2} \times \frac{\partial X_2}{\partial W_3} = -1 \times X_3 = -X_3$$




$$X'_2 \rightarrow W_3X_3 + W_4X_4$$

$$\frac{\partial f}{\partial W_3} = ?$$

$$\frac{\partial f}{\partial W_3} = \frac{\partial f}{\partial X_2} \times \frac{\partial X_2}{\partial X'_2} \times \frac{\partial X'_2}{\partial W_3} = -1 \times \sigma' \times X_3 = -\sigma'X_3$$

Example: Learning a Neural Network



Learning the Convolution filter by backpropagation:

Remember the convolution step:

X_{11}	X_{12}	X_{13}
X_{21}	X_{22}	X_{23}
X_{31}	X_{32}	X_{33}



Input

W_{11}	W_{12}
W_{21}	W_{22}

Filter

Output

Learning the Convolution filter by backpropagation:




(For detailed derivation check blackboard)

Learning the Convolution filter by backpropagation:

\mathbf{X}_{11}	\mathbf{X}_{12}	\mathbf{X}_{13}
\mathbf{X}_{21}	\mathbf{X}_{22}	\mathbf{X}_{23}
\mathbf{X}_{31}	\mathbf{X}_{32}	\mathbf{X}_{33}

Input



Filter

$\partial \mathbf{h}_{11}$	$\partial \mathbf{h}_{12}$
$\partial \mathbf{h}_{21}$	$\partial \mathbf{h}_{22}$


Error

(For detailed derivation check blackboard)

Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$





Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- Converted range: [0 1]


- Saturated activation causes gradients to vanish (over the chain of multiplication)
- Outputs are not zero centered
- Calculating “exponential” is computationally expensive.



Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$






Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



$\tanh(x)$


- Converted range: [-1 1]
- Zero Centered

- Saturated activation causes gradients to vanish (over the chain of multiplication)

Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$



ReLU


$$\max(0, x)$$



- $f(x) = \max(0, x)$
- Doesn't Saturate (+ve Region only)
- Converges faster than the previous two

Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




- Not zero centered
- Neurons with negative values will never activate!



Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$





Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




- $f(x) = \max(0.01x, x)$
- Doesn't Saturate
- Computationally efficient
- Converges faster than the previous two

- Not zero centered
- Neurons with negative values will never activate!

Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



$f(x)$

$$= \max(w_1^T x + b_1, w_2^T x + b_2)$$


- Doesn't saturate the activation.
- Generalizes ReLU and Leaky ReLU.

- Increased number of parameters!

Selection of Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




$$f(x)$$

$$= \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$

- Everything ReLU
- Closer to Zero mean
- Adds some robustness to noise.

- Calculating “exponential” is computationally expensive



Few more things

- Regularization
 - Normalization / Batch Normalization
 - Dropout
- Optimizers
 - SGD
 - SGD + Momentum
 - AdaGrad
 - RMSProp
 - Adam
- Hyper-parameter optimization
 - Grid Search
 - Bayesian Optimization
 - DoE
- Transfer Learning